

Common mistakes in



EF Core

NDC Porto 2023

Join the Conversation @SSW_TV @jernej_kavka #EFCore



Jernej Kavka (JK)

SSW Solution Architect



[@Jernej_kavka](https://twitter.com/Jernej_kavka)



github.com/jernejk



jkdev.me



linkedin.com/in/jernejkavka/



Brisbane Full Stack User Group



Host @ Global AI The Podcast

.NET Core and EF Core dev specialist

Microsoft AI MVP and ChatGPT enthusiast at night



Join the Conversation @SSW_TV @jernej_kavka #EFCore

Agenda



JAN

FEB

MAR

APR

MAY

MONTHLY TIP

SUNDAY

MONDAY

TUESDAY

PLAN SMART.

FILL OUT YOUR
WEEKLY AGENDA
NOTING WHEN YOU
WILL DO WHAT AND
HOW. THIS WILL
KEEP YOU ON TRACK
TO ATTAINING YOUR
GOAL.

Intro

NOTES

The 7 Deadly Sins

Mini Best Practices

We have lots of DB engines data access landscape



SQL Server



Cosmos DB



PostgreSQL



MySQL

Various ORM



EF Core (multiple DBs)



Dapper (multiple DBs)



Marten (PostgreSQL)



Cosmos DB SDK

Why EF Core

✓ Great balance

✦ Features

🔧 Flexibility

⚡ Performance

✓ Awesome for relational DBs

✓ Built-in migrations, DB scaffolding and other tooling

Maxim of quantity

be as informative as
required, but not more

Maxim of truth

only say things you know, or
reasonably believe, to be true

Maxim of relevance

don't include information you
do not know to be relevant

Maxim of clarity

avoid ambiguity and obscure language,
and present your message in a way
that is easy to process

Avoid the pain

- The EF Core issues that are solved
- Simplifying your code
- Benchmark your code with tests
- The horror I have seen... 🤖



Benchmark tools

- BenchmarkDotNet  **BenchmarkDotNet**
Powerful .NET library for benchmarking


Memory consumption and time difference

- Bombardier  **BOMBARDIER**

Simple CLI load testing tool

My testing environment

- .NET 7 with EF Core 7
- My beefy PC
 - CPU: Ryzen 9 5900X
 - Cores: 12
 - Threads: 24
 - Ram: 32GB



The screenshot displays the CPU-Z application window, which provides detailed information about the system's hardware. The 'CPU' tab is selected, showing the following specifications for the AMD Ryzen 9 5900X:

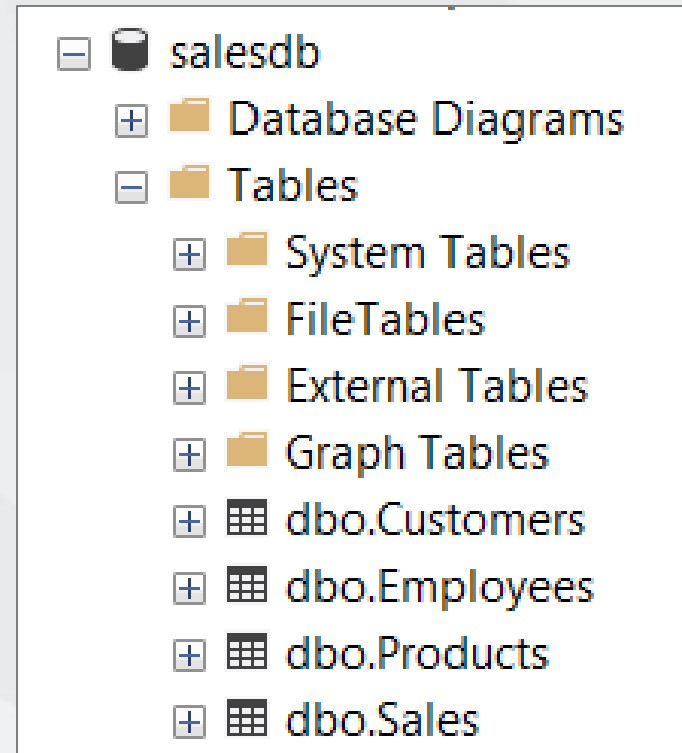
- Processor:**
 - Name: AMD Ryzen 9 5900X
 - Code Name: Vermeer
 - Package: Socket AM4 (1331)
 - Technology: 7 nm
 - Max TDP: 105.0 W
 - Core Voltage: 1.224 V
 - Specification: AMD Ryzen 9 5900X 12-Core Processor
 - Family: F
 - Ext. Family: 19
 - Model: 1
 - Ext. Model: 21
 - Stepping: 0
 - Revision: B0
 - Instructions: MMX(+), SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2, SSE4A, x86-64, AMD-V, AES, AVX, AVX2, FMA3, SHA
- Clocks (Core #0):**
 - Core Speed: 3592.72 MHz
 - Multiplier: x 36.0
 - Bus Speed: 99.80 MHz
 - Rated FSB: (empty)
- Cache:**
 - L1 Data: 12 x 32 KBytes (8-way)
 - L1 Inst.: 12 x 32 KBytes (8-way)
 - Level 2: 12 x 512 KBytes (8-way)
 - Level 3: 2 x 32 MBytes (16-way)
- Selection:** Socket #1
- Cores:** 12
- Threads:** 24

The bottom of the window shows the CPU-Z logo, version 1.97.0.x64, and buttons for Tools, Validate, and Close.

Test data



- 7 million sales
- 504 products
- 23 employees
- 20,000 customers



<https://www.sqlskills.com/sql-server-resources/sql-server-demos/>

A top-down view of a person's hands wearing black fingerless gloves, typing on a silver laptop keyboard. The laptop is open on a dark, textured surface. A semi-transparent dark grey rectangular box is centered over the image, containing the text 'The 7 Deadly Sins' followed by two icons: a red apple with a bite taken out of it, and a white chess knight piece.

The 7 Deadly Sins 🍏 🐎

The 7 Deadly EF Core Sins (easy to hard)

1. Casting **IQueryable** -> **IEnumerable**
2. Not using **AsNoTracking** (when appropriate)
3. Explicit joins
4. Getting all the columns
5. No pagination
6. Non-cancellable queries
7. Inefficient updates/deletes





"How bad can it be?"

(you never want the answer to this 🤝)

- Literally everyone before discovering the horrifying truth

Join the Conversation @SSW_TV @jerne_j_kavka #EFCore

#1 IQueryable -> IEnumerable

[Benchmark]

0 references | 0 changes | 0 authors, 0 changes

```
public int NaiveCount()
{
    return SalesDbContext.Sales
        .ToList()
        .Count;
}
```

[Benchmark]

0 references | 0 changes | 0 authors, 0 changes

```
public int QuerableToEnumerableCount()
{
    IEnumerable<Sale> sales = SalesDbContext.Sales;
    return sales.Count();
}
```

[Benchmark]

0 references | 0 changes | 0 authors, 0 changes

```
public int ImplicitEnumerableCount()
{
    return GetSalesEnumerable().Count();
}
```

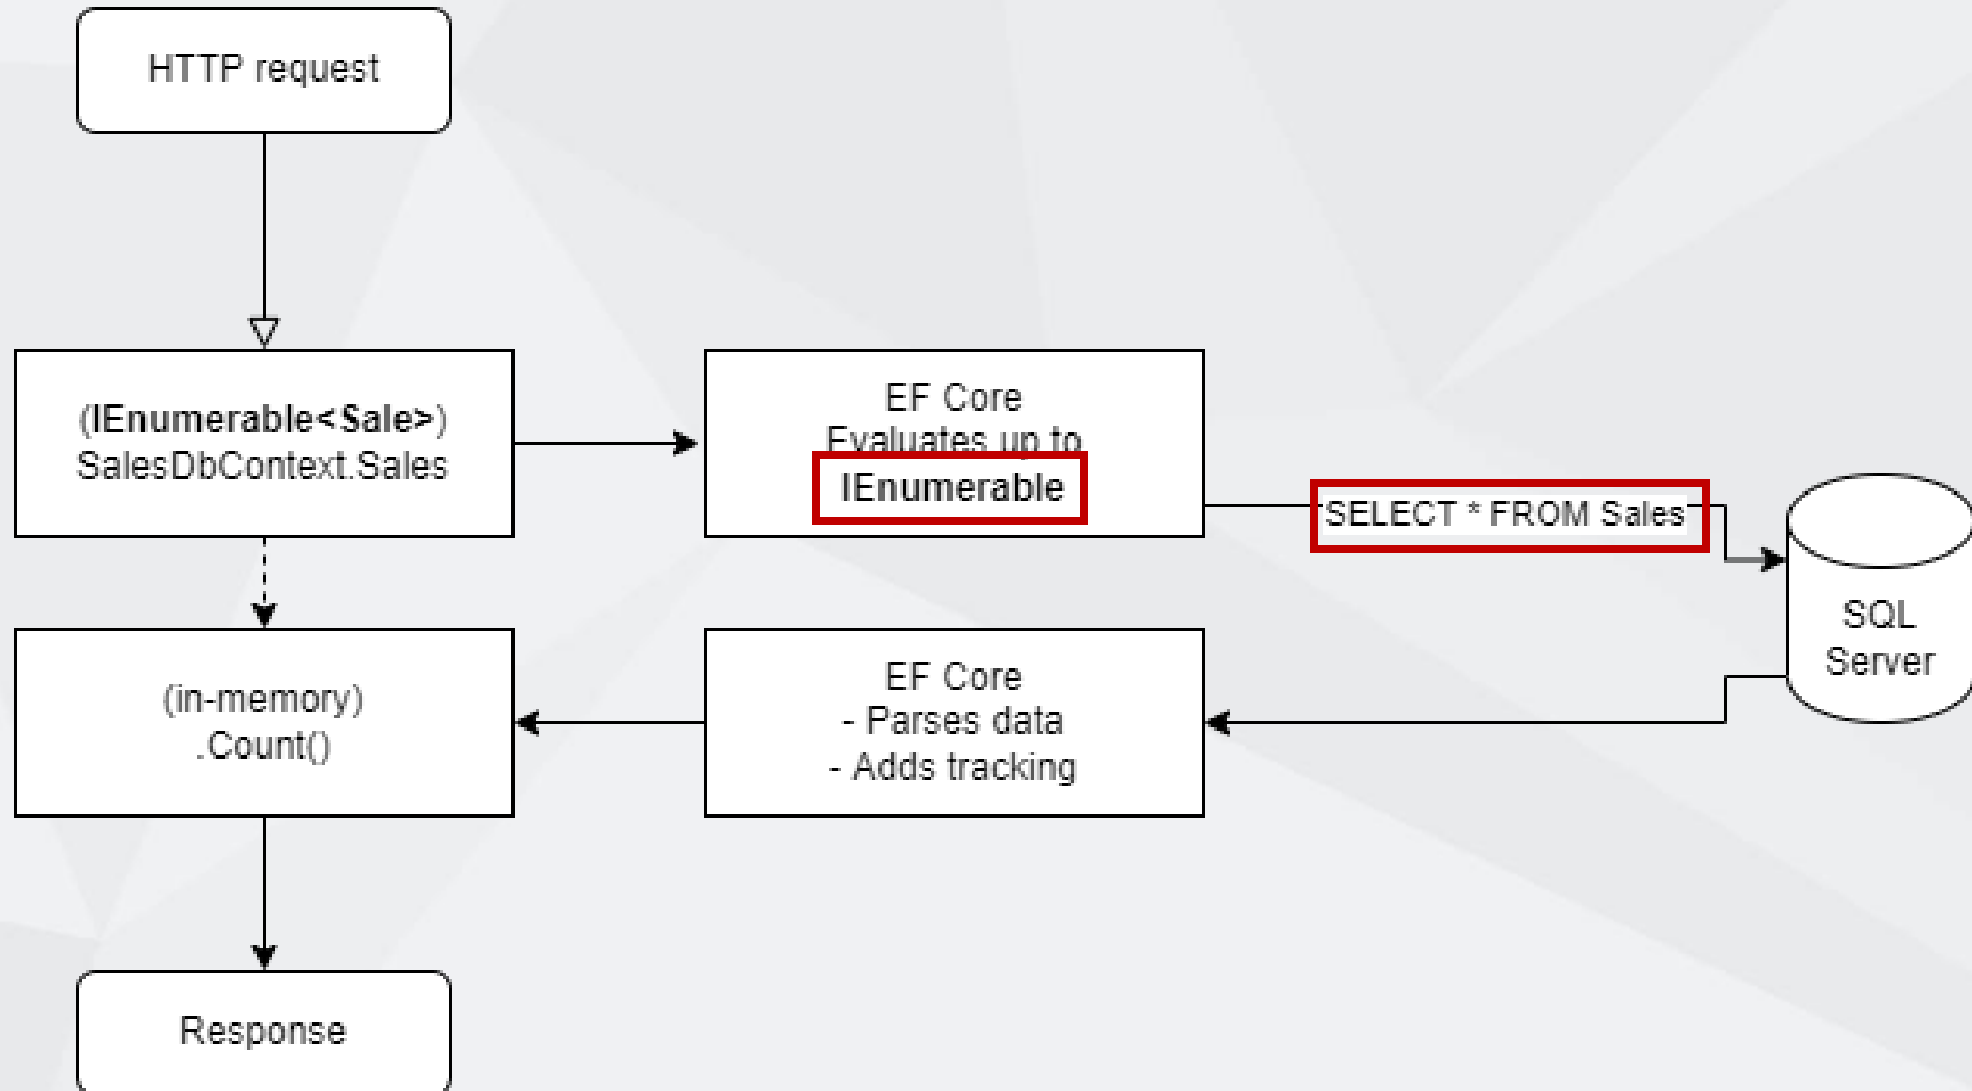
1 reference | 0 changes | 0 authors, 0 changes

```
private IEnumerable<Sale> GetSalesEnumerable()
{
    return SalesDbContext.Sales;
}
```



THEY ARE ALL
THE SAME

#1 What is going on?



#1 IQueryable

```
[Benchmark(Baseline = true)]
```

```
0 references | 0 changes | 0 authors, 0 changes
```

```
public int CountInDb()
```

```
{
```

```
    return GetSales().Count();
```

```
}
```

```
1 reference | 0 changes | 0 authors, 0 changes
```

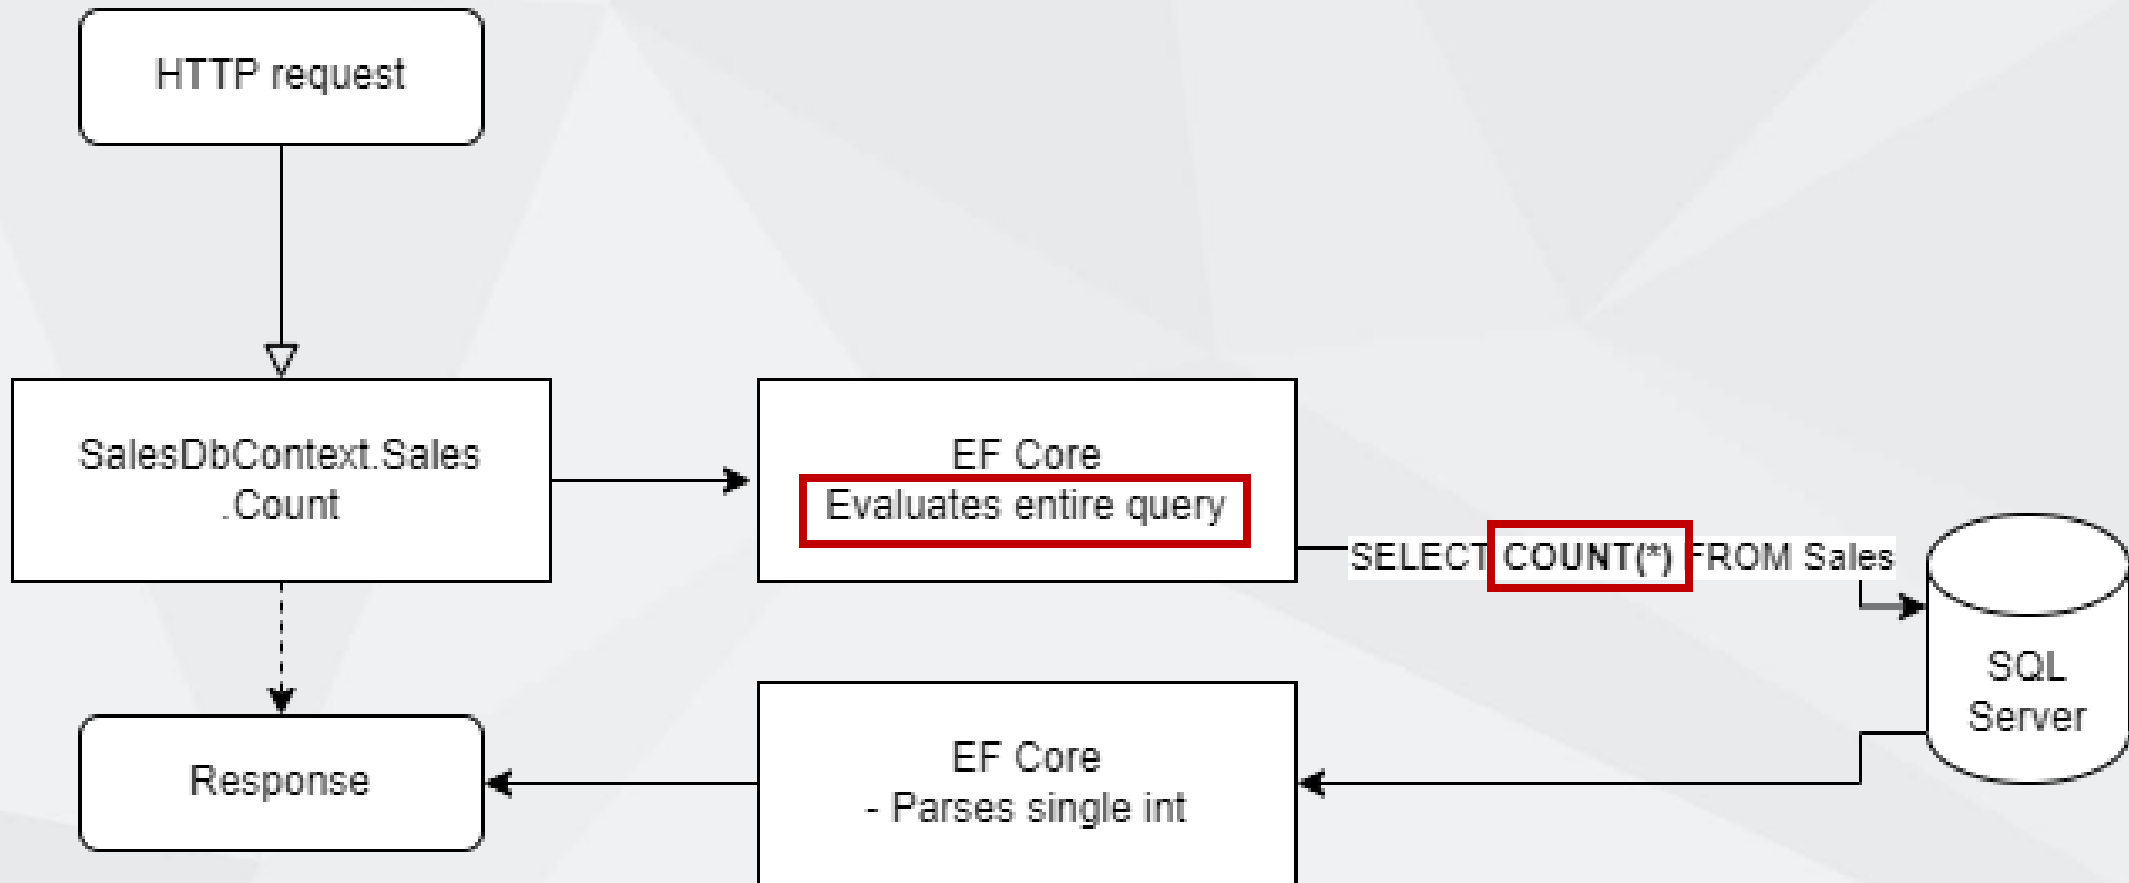
```
private IQueryable<Sale> GetSales()
```

```
{
```

```
    return SalesDbContext.Sales;
```

```
}
```

#1 How is this better?



IQueryable + EF Core (EF Core LINQ)
≠
IEnumerable (Standard LINQ)

#1 What is the memory impact?

Method	Mean	Error	StdDev	Ratio	RatioSD	Gen 0	Allocated
NaiveCount	3,675.44 ms	69.775 ms	71.654 ms	63.91	1.78	224000.0000	1,967,272 KB
CountInDb	57.52 ms	0.663 ms	0.588 ms	1.00	0.00	-	6 KB

Full 3.5 seconds faster!

64 times faster

~million times
less memory

#1 What about load testing?

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesCount/worstCase
Bombarding https://localhost:5001/ExamplesCount/worstCase for 1m0s using 5 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec        0.09       3.10    149.94
Latency      33.36s      7.62s    47.20s
```

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesCount/bestCase
Bombarding https://localhost:5001/ExamplesCount/bestCase for 1m0s using 5 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec       38.34      21.26    235.24
Latency     128.69ms     23.25ms    269.06ms
```

~260x faster on average in a load test!

#2 AsNoTracking

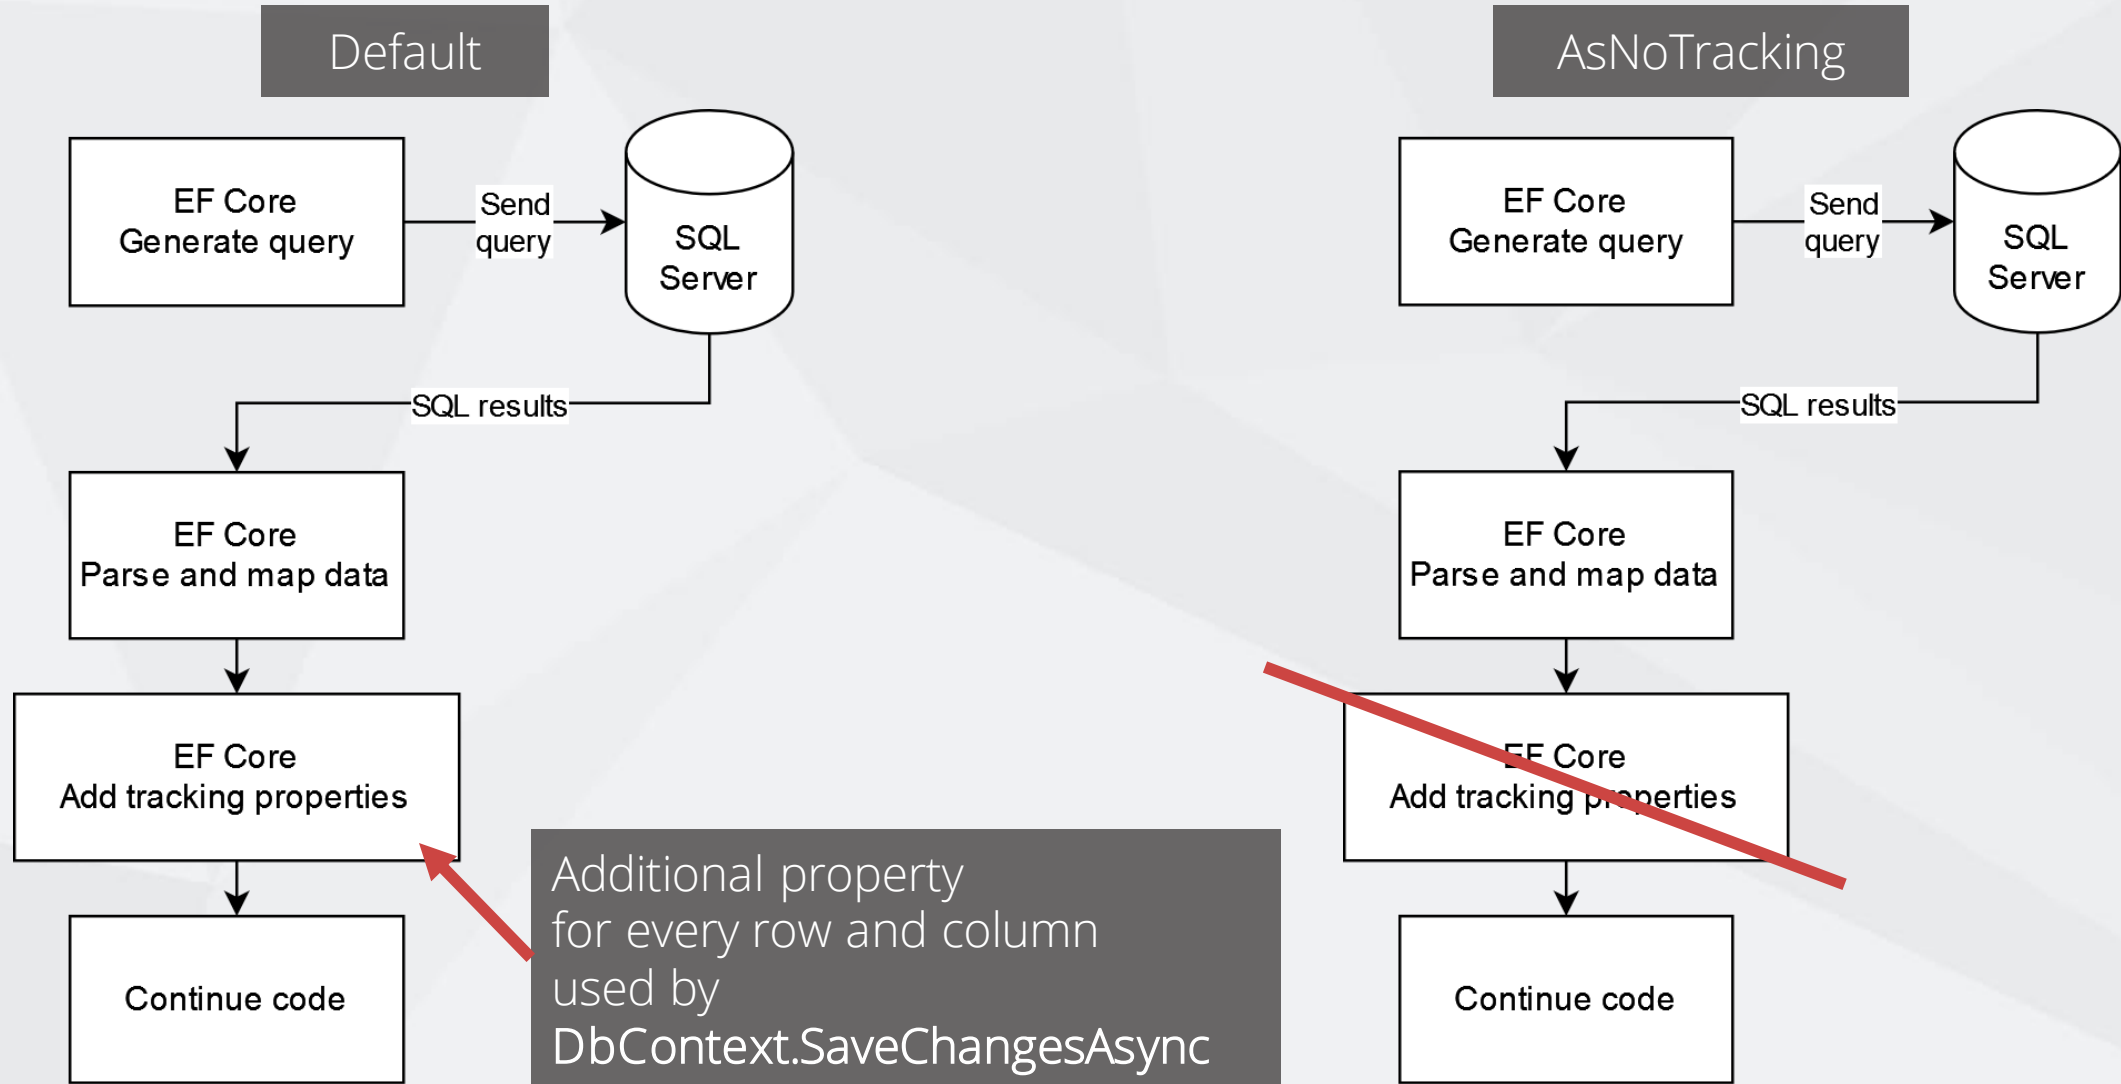
Exclude change tracking on entities

Less memory and CPU usage

⚠ Changes to the entities are **not tracked** ⚠

```
[Benchmark]
0 references | 0 changes | 0 authors, 0 changes
public List<Sale> GetSalesAsNoTracking()
{
    return SalesDbContext.Sales
        .AsNoTracking()
        .Where(x => x.CustomerId == 7)
        .ToList();
}
```

#2 What is happening?



#2 What is the impact?

Method	Mean	Error	StdDev	Ratio	Gen 0	Gen 1	Gen 2	Allocated
GetSalesQueryable	2,014.2 ms	39.37 ms	36.83 ms	1.00	65000.0000	19000.0000	2000.0000	582 MB
GetSalesAsNoTracking	387.4 ms	7.64 ms	13.58 ms	0.19	17000.0000	7000.0000	-	147 MB

~1.6 seconds faster!

~5 times faster

~4x less memory

⚠ Results be wildly different between queries and DB providers! ⚠
Some found it has no effect on SQLite DB provider

#2 What about load testing?

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesWhere/where
Bombarding https://localhost:5001/ExamplesWhere/where for 1m0s using 5 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec       3.68      12.10     199.96
Latency        1.40s     371.46ms    2.66s
```

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesWhere/whereNoTracking
Bombarding https://localhost:5001/ExamplesWhere/whereNoTracking for 1m0s using 5 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec      14.72      13.34     249.94
Latency       358.92ms    105.40ms    1.48s
```

~4x faster on average in a load test!

#3 Explicit includes

.Include(x => x.Customers)


They are always included even if not needed

Devs often forget to remove them

We get all columns even if we don't need them

#3 & #4 Implicit includes with Select

```
var query = _dbContext.Sales
    .AsNoTracking()
    .TagWithContext()
    .Include(x => x.SalesPerson)
    .Where(x => x.SalesPersonId == 1);
```



```
var query = _dbContext.Sales
    .AsNoTracking()
    .Where(x => x.SalesPersonId == 1)
    .Select(x => new SalesWithSalesPerson
    {
        CustomerId = x.CustomerId,
        SalesId = x.SalesPersonId,
        ProductId = x.ProductId,
        Quantity = x.Quantity,
        SalesPersonId = x.SalesPersonId,
        SalesPersonFirstName = x.SalesPerson.FirstName,
        SalesPersonLastName = x.SalesPerson.LastName
    });
```

#3 & #4 Load Test

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesJoin/badCase
Bombarding https://localhost:5001/ExamplesJoin/badCase for 1m0s using 5 connection(s)
```

```
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec        2.90       11.85    238.05
Latency         1.92s     561.85ms    5.16s
```

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesJoin/implicitJoin
Bombarding https://localhost:5001/ExamplesJoin/implicitJoin for 1m0s using 5 connection(s)
```

```
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec        7.49       25.53    263.10
Latency        748.18ms    575.21ms    5.30s
```

~2.5x faster on average in a load test!

#5 Common Pagination issues

- Inefficient code
 - Filter data
 - Get **all** filtered data
 - Count them
 - Apply Take and Skip **in-memory** for pagination

#5 Pagination code – naïve

```
IQueryable<Sale> query = _dbContext.Sales
    .AsNoTracking()
    // Very commonly used in combination of in-memory pagination, making things even worse
    .Include(x => x.SalesPerson)
    .Where(x => x.SalesPersonId == salesPersonId);

List<Sale> dbResult = await query.ToListAsync(ct);
int count = dbResult.Count;
List<SalesWithSalesPerson> result = dbResult
    .Skip(page * pageSize)
    .Take(pageSize)
    .Select(x => new SalesWithSalesPerson
    {
        CustomerId = x.CustomerId,
        SalesId = x.SalesPersonId,
        ProductId = x.ProductId,
        Quantity = x.Quantity,
        SalesPersonId = x.SalesPersonId,
        SalesPersonFirstName = x.SalesPerson.FirstName,
        SalesPersonLastName = x.SalesPerson.LastName
    })
    .ToList();
```

Gets all filtered data

Count in-memory

Pagination is done
In-memory

#5 Pagination code – improved

```
IQueryable<SalesWithSalesPerson> query = _dbContext.Sales
    .AsNoTracking()
    .Where(x => x.SalesPersonId == salesPersonId)
    .Select(x => new SalesWithSalesPerson
    {
        CustomerId = x.CustomerId,
        SalesId = x.SalesPersonId,
        ProductId = x.ProductId,
        Quantity = x.Quantity,
        SalesPersonId = x.SalesPersonId,
        SalesPersonFirstName = x.SalesPerson.FirstName,
        SalesPersonLastName = x.SalesPerson.LastName
    });
```

We are preparing the query
For counting and pagination

```
// After all conditions are applied, count them in DB.
```

```
int count = await query.CountAsync(ct)
```

Count in SQL Server

```
// Apply paginations, sorts and more complex select statements.
```

```
query = query
```

```
.Skip(page * pageSize)
.Take(pageSize);
```

Pagination is done
on SQL Server

```
List<SalesWithSalesPerson> result = await query.ToListAsync(ct)
```

#5 Load test

```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesPaginations/worstCase
Bombarding https://localhost:5001/ExamplesPaginations/worstCase for 1m0s using 5 connection(s)
[=====
Done!
Statistics      Avg      Stdev      Max
Reqs/sec       3.72      17.21     210.48
Latency        1.35s     365.60ms    2.27s
```

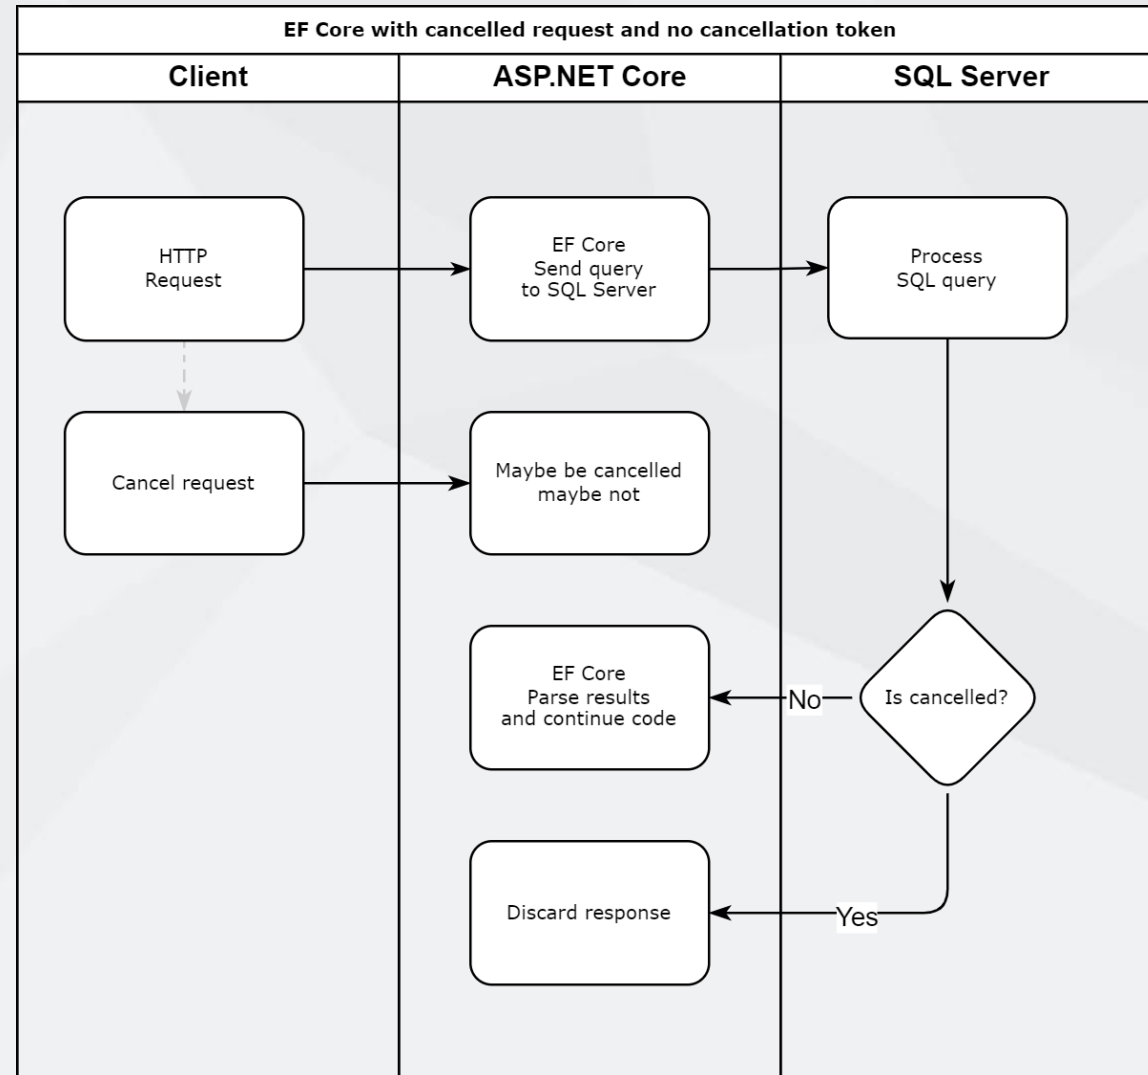
```
> bombardier -c 5 -t 60s -d 60s -l https://localhost:5001/ExamplesPaginations/executedOnDB
Bombarding https://localhost:5001/ExamplesPaginations/executedOnDB for 1m0s using 5 connection(s)
[=====
Done!
Statistics      Avg      Stdev      Max
Reqs/sec      65.41      21.98     263.10
Latency       76.94ms     13.79ms    237.05ms
```

~18x faster on average in a load test!

#6 Non-cancellable queries

- When query is running on SQL Server
 - It will run until completion
- Complex queries can stuff up SQL Server

#6 What's the problem?




#6 Async and CancellationToken

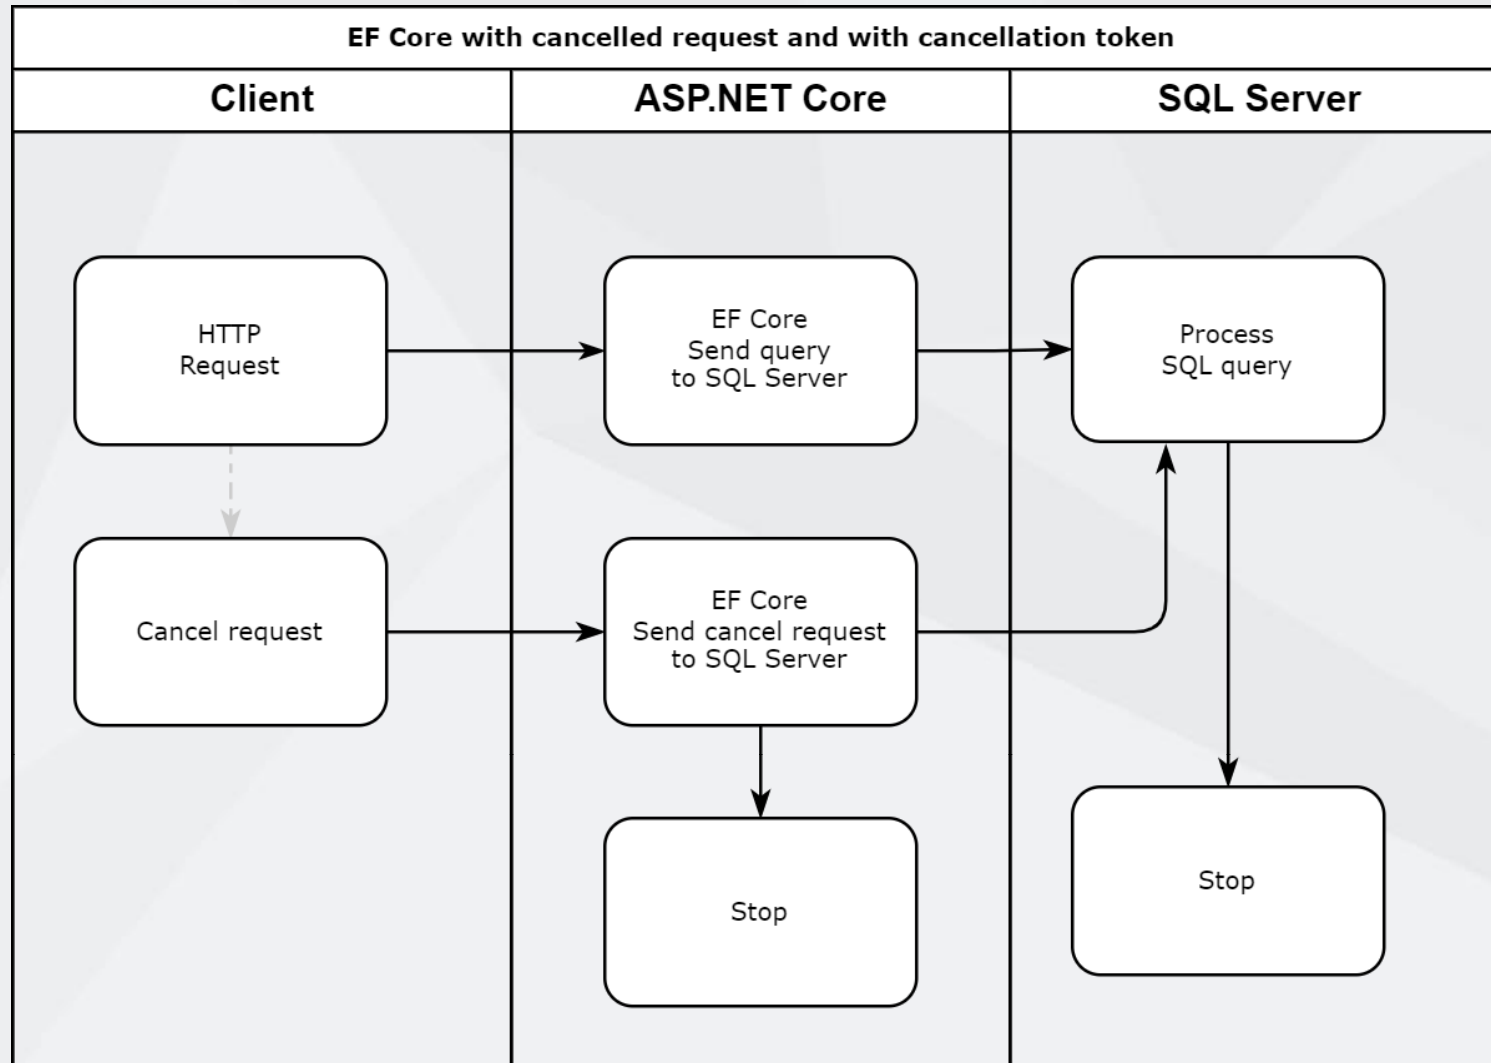
Stop execution on the server when cancelled

```
[HttpGet("splitQuery")]
0 references | Jernej Kavka (JK), 65 days ago | 1 author, 1 change
public async Task<TestResult<int>> SplitQuery(CancellationToken ct)
{
    var query = _dbContext.Sales
        .AsNoTracking()
        .AsSplitQuery()
        .Where(x => x.SalesPersonId == 1)
        .Select(x => new SalesWithSalesPerson
        {
            CustomerId = x.CustomerId,
            SalesId = x.SalesPersonId,
            ProductId = x.ProductId,
            Quantity = x.Quantity,
            SalesPersonId = x.SalesPersonId,
            SalesPersonFirstName = x.SalesPerson.FirstName,
            SalesPersonLastName = x.SalesPerson.LastName
        });

    List<SalesWithSalesPerson> result = await query.TagWithContext().ToListAsync(ct);
}
```



#6 How it works cancellation token?



#6 Load test

```
> bombardier -c 12 -t 1s -d 60s -l https://localhost:5001/test/count
Bombarding https://localhost:5001/test/count for 1m0s using 12 connection(s)
[=====
Done!
Statistics      Avg      Stdev      Max
Reqs/sec       1.64      28.10     666.49
Latency        10.14s     2.83s     17.89s
```

```
> bombardier -c 12 -t 1s -d 60s -l https://localhost:5001/test/countasync-ct
Bombarding https://localhost:5001/test/countasync-ct for 1m0s using 12 connection(s)
[=====
Done!
Statistics      Avg      Stdev      Max
Reqs/sec       2.04      32.15     631.45
Latency         5.85s    368.77ms     7.26s
```

~2x faster on average in a load test!
Also, responsiveness returns 40 seconds faster

#7 Inefficient update/delete #1

- Before EF7 we had to fetch data before could update
- Entities need to be tracking
- Can be very inefficient

#7 Inefficient update/delete #2

```
IQueryable<Employee> query = GetBaseQuery(isLoadFriendly);
var employees = query.ToList();
foreach (var employee in employees)
{
    string firstName = employee.FirstName;
    employee.FirstName = employee.LastName;
    employee.LastName = firstName;
}

_dbContext.SaveChangesAsync();
```

-- Get all employees that we want to swap first and last name.
-- We need to fetch all columns as it needs to be tracking!

```
SELECT TOP(@__p_0) [e].[EmployeeID], [e].[FirstName], [e].[LastName], [e].[MiddleInitial]
FROM [Employees] AS [e]
```

```
SET NOCOUNT ON;
UPDATE [Employees] SET [FirstName] = @p0, [LastName] = @p1
OUTPUT 1
WHERE [EmployeeID] = @p2;
```

-- Lots of updates...

```
UPDATE [Employees] SET [FirstName] = @p63, [LastName] = @p64
OUTPUT 1
WHERE [EmployeeID] = @p65;
UPDATE [Employees] SET [FirstName] = @p66, [LastName] = @p67
OUTPUT 1
WHERE [EmployeeID] = @p68;
```

#7 Inefficient update/delete #3

```
IQueryable<Employee> query = GetBaseQuery(isLoadFriendly);  
var employees = query  
    .TagWith("Swap first and last name")  
    .ExecuteUpdate(x => x  
        .SetProperty(p => p.FirstName, b => b.LastName)  
        .SetProperty(p => p.LastName, b => b.FirstName));  
  
_dbContext.SaveChangesAsync();
```

-- Swap first and last name

```
UPDATE [e]  
SET [e].[LastName] = [e].[FirstName],  
    [e].[FirstName] = [e].[LastName]  
FROM [Employees] AS [e]
```

#7 Inefficient update/delete #4

```
~#@> bombardier -c 12 -t 1s -d 60s -l https://localhost:5001/ExamplesUpdate/worstCase
Bombarding https://localhost:5001/ExamplesUpdate/worstCase for 1m0s using 12 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec      524.96    478.57    3168.91
Latency       22.87ms   54.64ms   1.28s
```

```
~#@> bombardier -c 12 -t 1s -d 60s -l https://localhost:5001/ExamplesUpdate/updateQuery
Bombarding https://localhost:5001/ExamplesUpdate/updateQuery for 1m0s using 12 connection(s)
[=====]
Done!
Statistics      Avg      Stdev      Max
Reqs/sec     1659.20    479.76    2781.09
Latency       7.23ms    26.63ms   1.66s
```

For only 23 records, it's still ~3x faster!
Performance improvements can be massive depending on the query!

Mini Best Practices

Bonus #1 - Other things I have seen

0 references | 0 changes | 0 authors, 0 changes

```
public bool HasAny()
{
    var employees = _dbContext.Employees
        .Where(x => x.EmployeeId == 1)
        .ToList();
    return employees.Count > 0;
}
```

0 references | 0 changes | 0 authors, 0 changes

```
public bool HasAny()
{
    return _dbContext.Employees
        .Any(x => x.EmployeeId == 1);
}
```

0 references | 0 changes | 0 authors, 0 changes

```
public List<SaleModel> GetSales(int? saleId)
{
    var sales = _dbContext.Sales
        .Include(x => x.Product)
        .ToList();

    if (saleId != null)
    {
        sales = sales
            .Where(x => x.SalesId == saleId.Value)
            .ToList();
    }

    return sales.Select(x => new SaleModel
    {
        SaleId = x.SalesId,
        Quantity = x.Quantity,
        ProductName = x.Product.Name
    }).ToList();
}
```

0 references | 0 changes | 0 authors, 0 changes

```
public List<SaleModel> GetSales(int? saleId)
{
    IQueryable<Sale> salesQuery = _dbContext.Sales;

    if (saleId != null)
    {
        salesQuery = salesQuery
            .Where(x => x.SalesId == saleId.Value);
    }

    return salesQuery.Select(x => new SaleModel
    {
        SaleId = x.SalesId,
        Quantity = x.Quantity,
        ProductName = x.Product.Name
    }).ToList();
}
```


Bonus #2

DbContextPool

- Reuses existing **DbContext**
- Fewer connections to SQL Server
- Improves resiliency

```
services.AddDbContextPool<SalesDbContext>(options =>  
{  
    options.UseSqlServer(Configuration.GetConnectionString("SalesDB"));  
});
```

Bonus #3

TagWith

Adds comments into SQL query for easier debugging

```
var list = await _context.Tweets
    .TagWith("GetTweets")
    .ToListAsync(cancellationToken)
    .ConfigureAwait(false);
```



```
-- GetTweets
```

```
SELECT [t].[Id], [t].[CreatedUtc], [t].[Message], [t].[Username]
FROM [Tweets] AS [t]
```

Bonus #4 – ChatGPT

ChatGPT **PLUS**

You're an EF Core expert and clean code developer.

<user's instructions>

DbContext:

<Optional project's DbContext, can be trimmed to include only the relevant parts>

Entities:

<Only necessary entities>

Return types:

<Optional return type we want to map>



ChatGPT may produce inaccurate information about people, places, or facts. [ChatGPT May 12 Version](#)

Join the Conversation @SSW_TV @jernej_kavka #EFCore

Bonus #4 – ChatGPT

🌟 Model: GPT-4



You're an EF Core expert and clean code developer.

Get all employees details, for each get statistics on how many products they earned and a list of sales they did, all paginated.

DbContext:

```
public partial class SalesDbContext : DbContext
{
    public SalesDbContext(DbContextOptions<SalesDbContext> options)
        : base(options)
    {
    }

    public virtual DbSet<Customer> Customers { get; set; }
```

```
public async Task<List<EmployerStats>> GetEmployeeStats(int pageNumber,
{
    return await _context.Employees
        .OrderBy(e => e.EmployeeId)
        .Skip((pageNumber - 1) * pageSize)
        .Take(pageSize)
        .Select(e => new EmployerStats
        {
            FirstName = e.FirstName,
            LastName = e.LastName,
            TotalProductSold = e.Sales.Count,
            TotalSales = e.Sales.Sum(s => s.Quantity * s.Product.Price),
            TotalQuantity = e.Sales.Sum(s => s.Quantity),
            Sales = e.Sales.Select(s => new SaleModel
            {
                SaleId = s.SalesId,
                ProductName = s.Product.Name,
                Quantity = s.Quantity
            })
        })
        .ToListAsync();
}
```

Before we end #1

When is something not premature optimization?

- ✓ Following best practices for given scenario

- ✓ When high-load is expected

 - ⓘ Talk to Product Owner and make it part of acceptance criteria

- ✓ When it already causes performance issues

 - ⓘ Talk to Product Owner and make it part of acceptance criteria or create tech debt PBI/issue and mention it in the code

 - 🌐 ssw.com.au/rules/technical-debt/

Before we end #2

What should you not do?

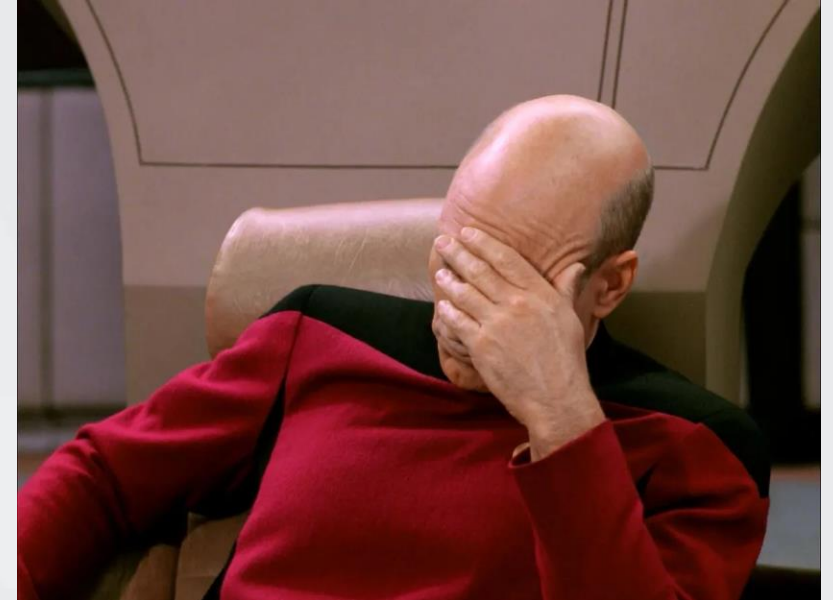
✗ “We’ll fix it later” with no actions

ℹ Create a tech debt PBI/issue and a link to it in code as a TODO comment

🌐 ssw.com.au/rules/technical-debt/

✗ “Premature optimization” as an excuse for bad code

ℹ There are valid reason to not optimize your code, bad code is not one of them!



Resources

Source code for examples and benchmarking

github.com/jernejk/EFCoreSamples.StabilityAndPerformance

Microsoft resources for better EF Core performance

docs.microsoft.com/en-us/ef/core/performance/

Rules to Better Entity Framework

ssw.com.au/rules/rules-to-better-entity-framework

Rules to Better LINQ

ssw.com.au/rules/rules-to-better-linq

Nick Chapsas YouTube – Great breakdowns of features and best practices

youtube.com/watch?v=Q4LtKa_HTHU

Join the Conversation @SSW_TV @jernejkavka #EFCore

Key takeaways 🗝️

- **IEnumerable** is dangerous when used with EF Core
- **AsNoTracking** has massive performance impact
 - Make sure you don't need to update/delete the entities
- Simpler code can result in better performance

Find this presentation on GitHub!



github.com/sswconsulting/presentations

The background is a light gray with a subtle geometric pattern of overlapping triangles. On the left side, there are colorful streamers in red, blue, and purple, along with various shapes like triangles and squares in matching colors. On the right side, there are stylized fireworks in red and blue, exploding upwards. The overall theme is celebratory and festive.

Thank you!

info@ssw.com.au

www.ssw.com.au

Sydney | Melbourne | Brisbane

Questions?

info@ssw.com.au

www.ssw.com.au

Sydney | Melbourne | Brisbane

